

Parallel Programming

Parallel Algorithm Design

Jan Schönherr

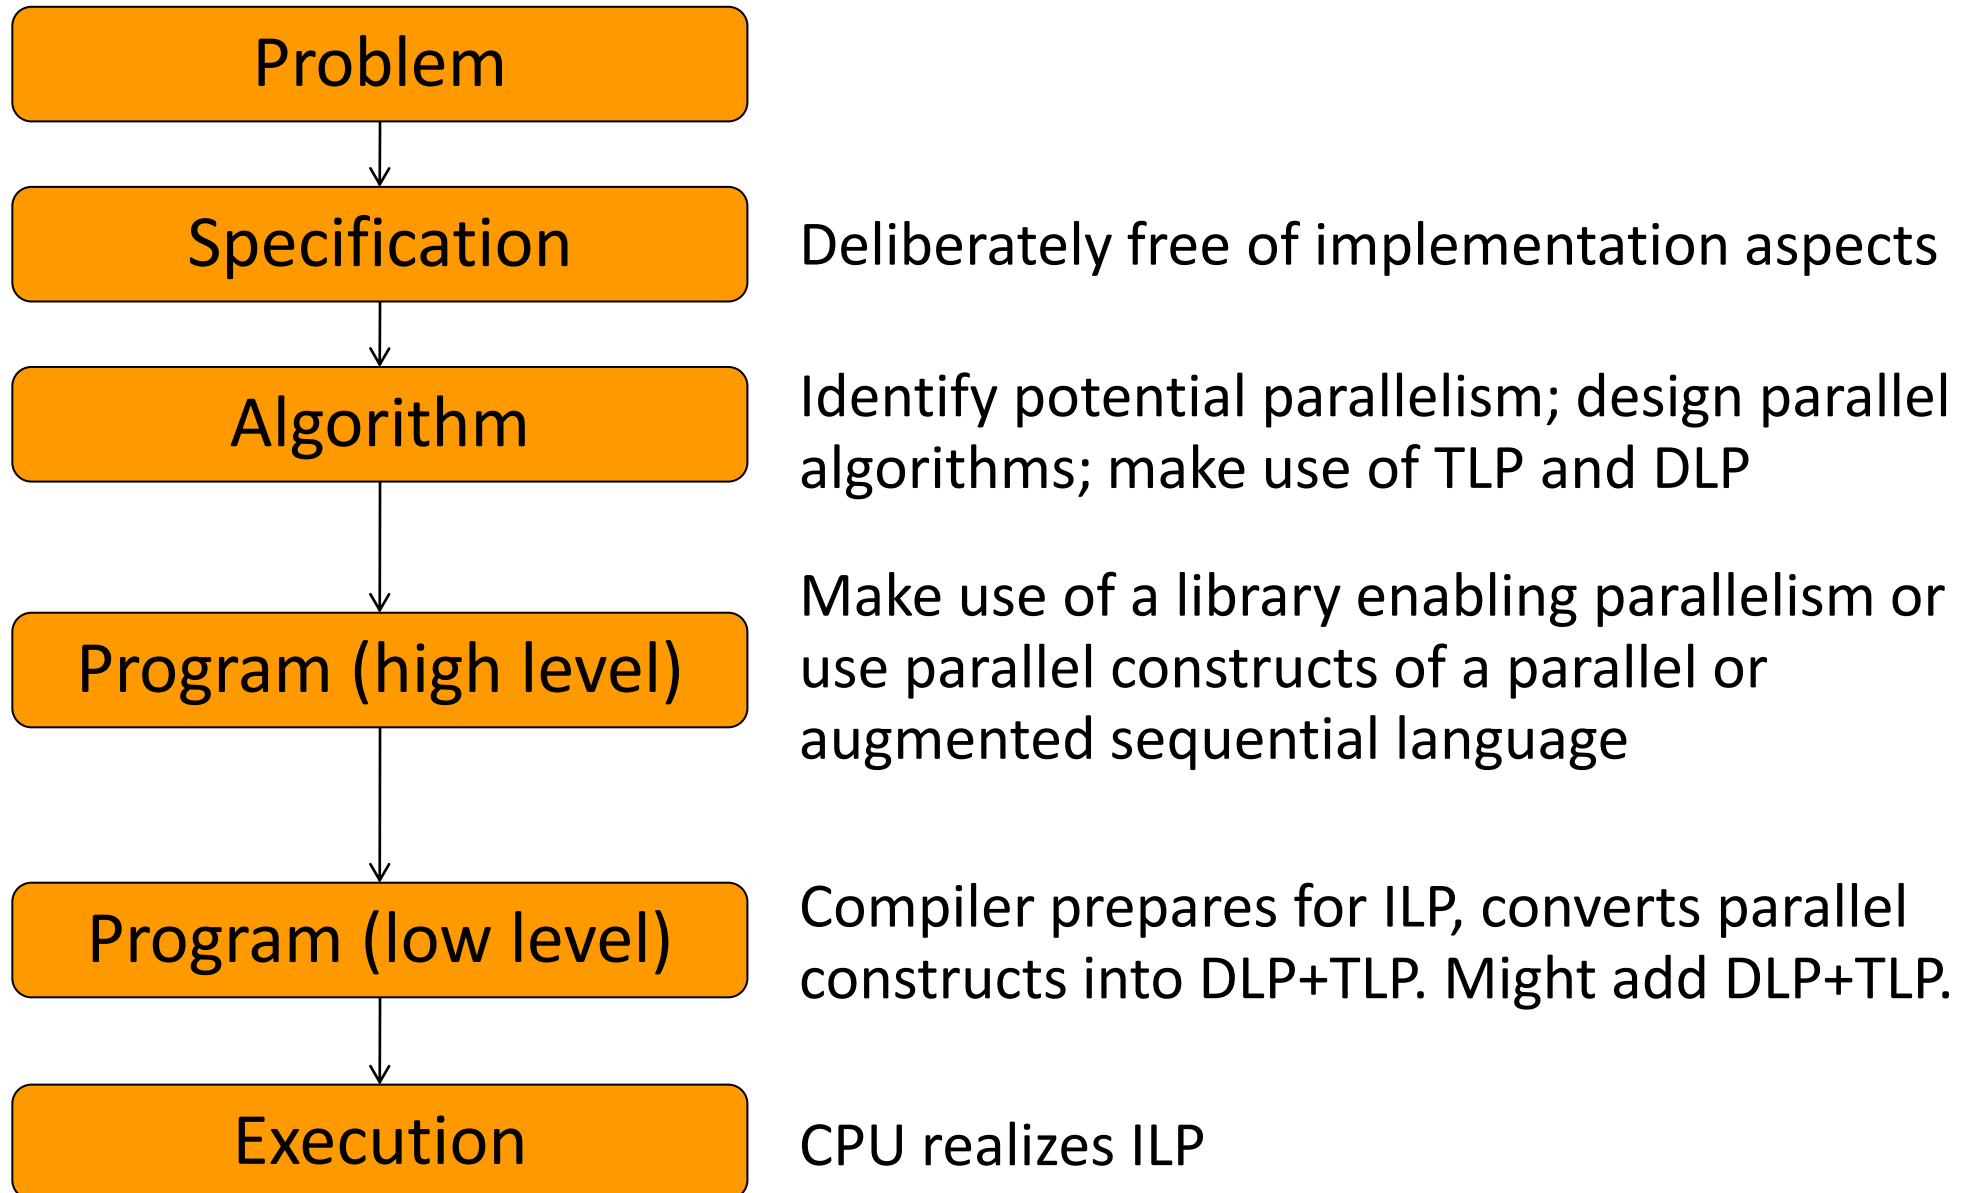
Technische Universität Berlin

School IV – Electrical Engineering and Computer Sciences

Communication and Operating Systems (KBS)

Einsteinufer 17, Sekr. EN6, 10587 Berlin

Parallelism in Software Development



Foster's Design Methodology

- > Encourages scalable parallel algorithms
 - > Targets mainly distributed memory system
 - > Applicable to other systems as well
- > Published 1995
- > Four steps

<ol style="list-style-type: none"> 1. Partitioning 2. Communication 	}	problem specific
<ol style="list-style-type: none"> 3. Agglomeration 4. Mapping 	}	system specific

(available online: <http://www.mcs.anl.gov/~itf/dbpp>)

Partitioning

- > Fine-grained decomposition into primitive tasks
 - > Task (or functional) decomposition
 - > Decompose based on central computation
 - > What data is needed to perform a task?
 - > Data (or domain) decomposition
 - > Decompose based on central data structures
 - > How can operations on decomposed data structures be realized?
- > Always consider both and different variants of them
- > Decomposition gives upper bound on parallelism
- > Goals
 - > A lot more primitive tasks than processing/memory elements
 - > Number of tasks scales with the problem size

Communication

- > Dependency analysis between primitive tasks
 - > Data or temporal dependencies
- > Typical properties
 - > Local or global dependencies
 - > i. e., communicate with few or many
 - > Structured or unstructured
 - > Structures can be taken advantage of later
 - > Static or dynamic communication partners
 - > Synchronous or asynchronous
 - > Do we now beforehand what data is needed where?
- > Goals for communication patterns
 - > Local, balanced, concurrent

Agglomeration

- > Group primitive tasks into larger tasks
 - > Reduce dependencies and overhead, increase locality and concurrency
 - > Trade off computation and communication
 - > Replication of data and/or computation?
- > Goals
 - > Structure algorithm towards real system
 - > Retain scalability
 - > Balance computation and communication (if beneficial)
 - > Especially in one-task-per-processor situations

Mapping

- > Map agglomerated tasks to processing/memory elements
 - > Externally: Create one process/thread per task
 - > Mapping handled by the operating system
 - > Internally: Handle multiple tasks per thread/process

- > Goals
 - > Minimized execution time
 - > Overlap communication with computation
 - > Load balancing

Common Algorithm Structures

- > Organized by tasks
 - > Task parallelism
 - > Divide and Conquer
- > Organized by data decomposition
 - > Geometric decomposition
 - > Recursive data
- > Organized by data flow
 - > Pipeline
 - > Event-based coordination

Task parallelism

- > Many tasks allow balanced scheduling
 - > Useful when tasks complexity varies or the system is heterogeneous in some way
 - > Rule of thumb: have at least 10 times more tasks than processing elements (but still consider management overhead)
 - > Task scheduling
 - > Static vs. dynamic
 - > Centralized vs. decentralized work queue

- > Embarrassingly parallel tasks
 - > No dependencies between tasks

Divide and Conquer

- > Principle
 - > Split problem into subproblems recursively
 - > Solve primitive subproblems
 - > Combine partial solutions to one solution recursively
- > Recursive generation of tasks
 - > Fall back to sequential solution at some point
- > Does not fully utilize system at the beginning and the end
 - > Divide into more than two subproblems
 - > Combine more than two partial solutions
 - > Divide/combine in parallel

Geometric Decomposition

- > Split a regular data structure into regular chunks
 - > Blocked and cyclic layouts along one or more dimensions

1D block

P0
P1
P2
P3

2D block

P0	P1
P2	P3

2D block-cyclic

P0	P1	P0	P1
P2	P3	P2	P3
P0	P1	P0	P1
P2	P3	P2	P3


- > Cyclic layouts enable static load balancing
- > Dynamic assignment of chunks also possible

Recursive Data

- > Seemingly sequential operations on recursive data structures
 - > e. g., linked lists, trees, graphs, ...
- > Divide & Conquer works in some cases
- > Rethink the problem in terms of operations on every member (e. g. recursive doubling)
 - > Might increase algorithmic complexity, but could still be faster for a certain number of processors
 - > Might work well on SIMD architectures

Pipeline

- > For simple ordering constraints
- > Throughput oriented



- > Concurrency limited by the number of stages
- > Slowest stage might become a bottleneck
 - > Parallelize stages
 - > Have multiple pipelines

Event-Based Coordination

- > For irregular, dynamic ordering constraints
 - > “Generalized pipeline”
- > Tasks wait for events, process them and generate events themselves
- > Requires one or multiple event queues

More complex problems

- > Break it down into manageable pieces
 - > Express problem in terms of known problems
 - > Apply existing algorithms
 - > Design missing pieces (e. g., with Foster's methodology)

- > Benefit from work of others
 - > Libraries with already parallelized primitives
 - > Optimized and improved over (a long) time
 - > Remaining task: reorganize data between function calls