

Parallel Programming

Coprocessors

Jan Schönherr

Technische Universität Berlin

School IV – Electrical Engineering and Computer Sciences

Communication and Operating Systems (KBS)

Einsteinufer 17, Sekr. EN6, 10587 Berlin

Coprocessor

- > A coprocessor is an additional (usually optional) processing unit, which is fed work by the CPU.
 - > Frees CPU resources for other activities
 - > Usually specialized, i. e., faster, better energy efficiency

- > For example
 - > Intel 80x87
 - > SPEs in the Cell processor
 - > GPUs (also integrated GPUs)
 - > Intel Xeon Phi

CPU vs. Coprocessor

- > General Purpose CPU
 - > Not bad in most use cases, not good in most use cases
- > Specialized Coprocessor
 - > Particular good in its specific use case
 - > Particular bad in everything else
- > Chicken/egg problem
 - > CPU is designed to execute existing software faster
 - > Software is written for existing CPUs

Examples

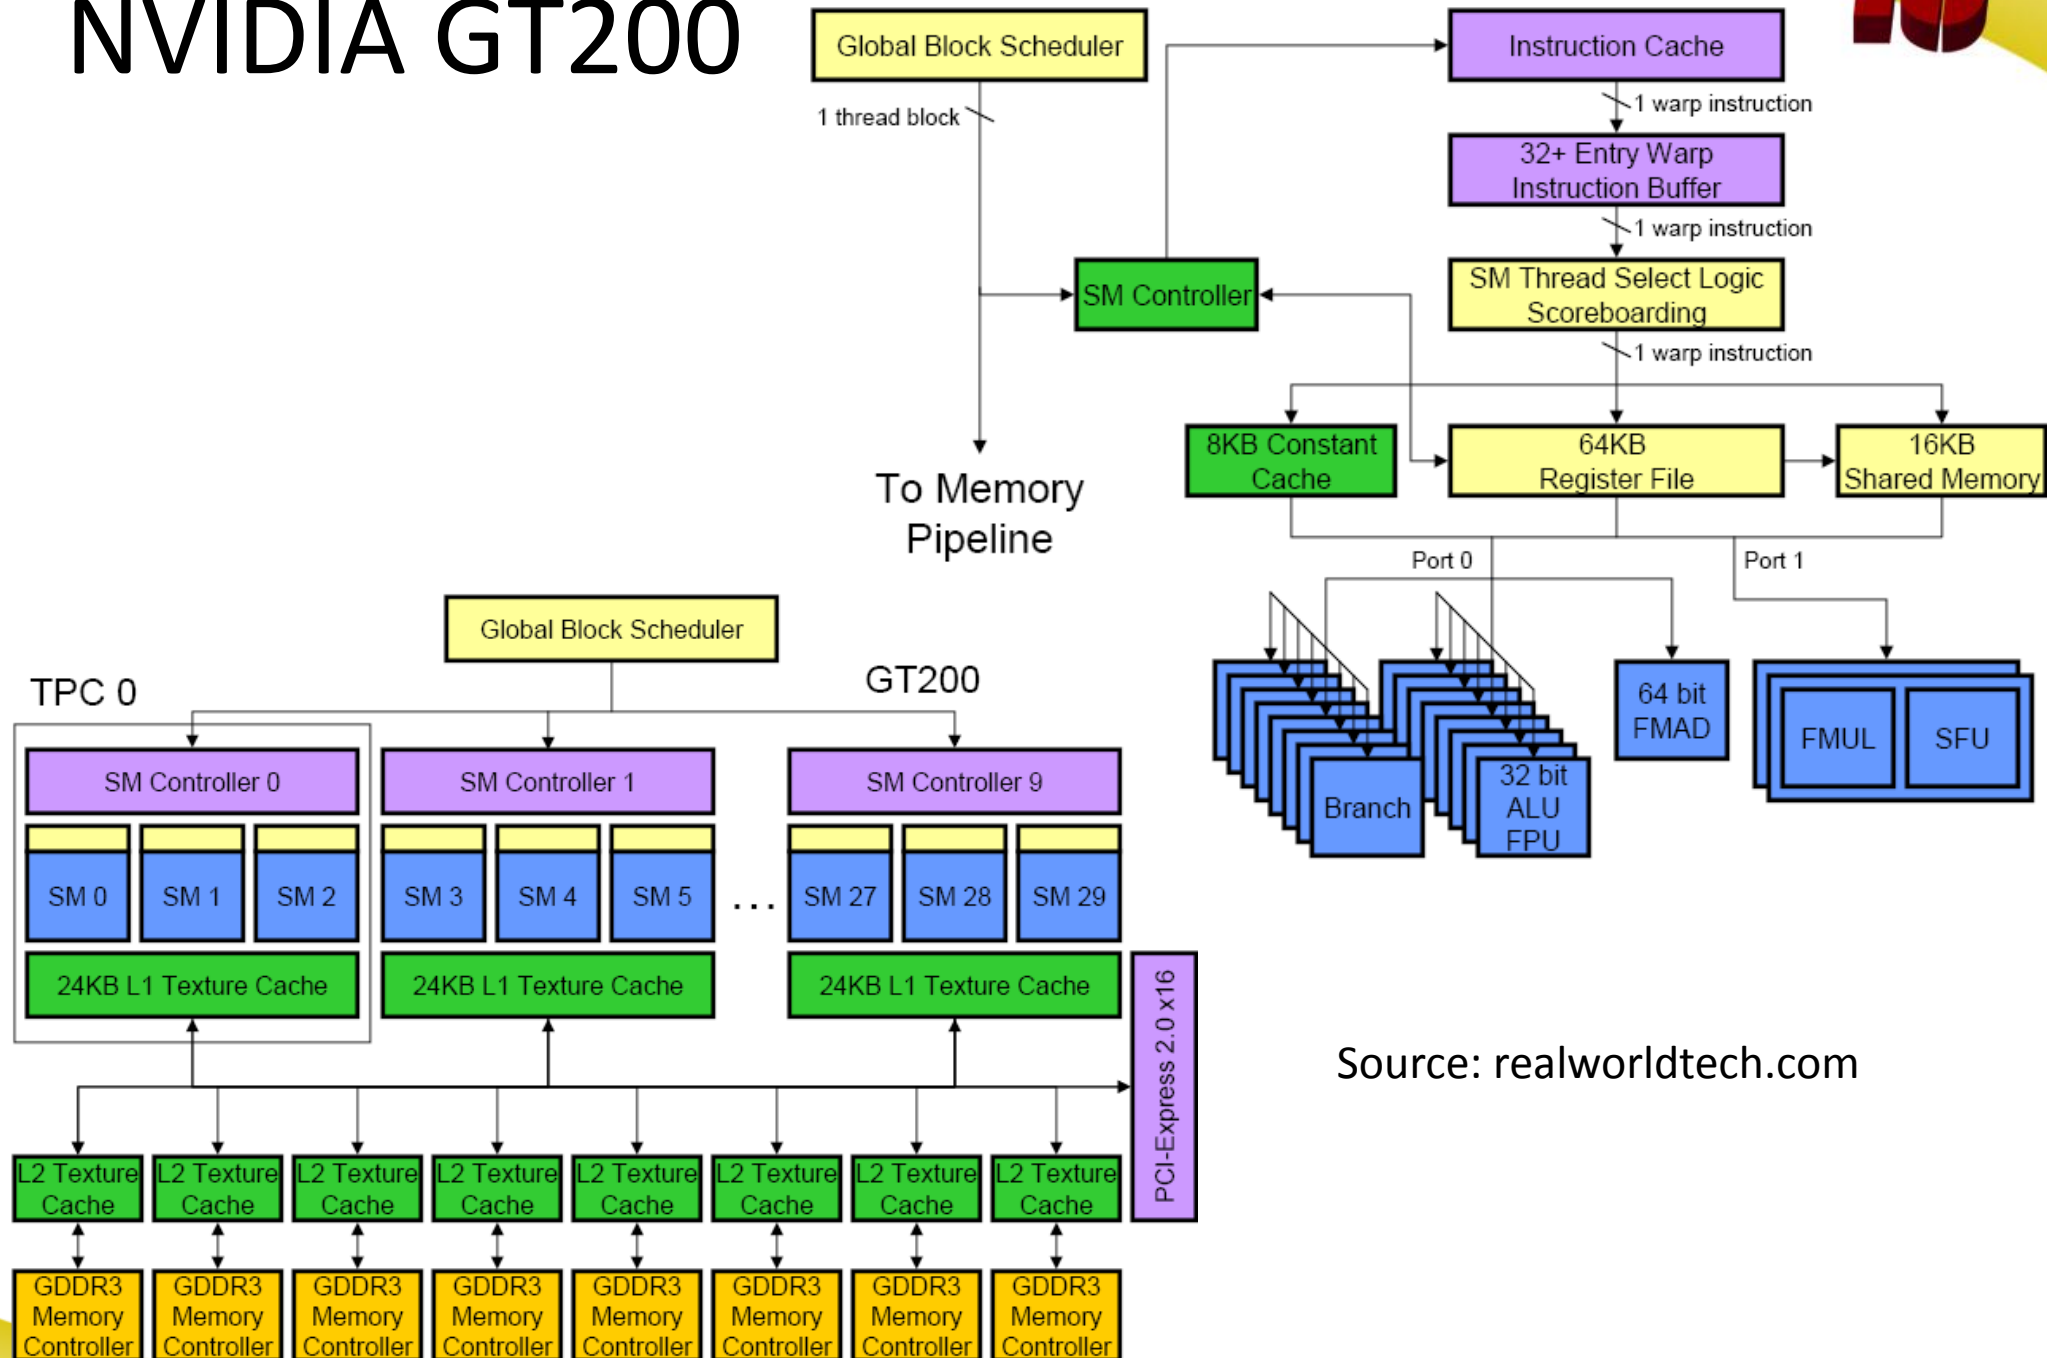
- > PowerXCell 8i processor
 - > 1 PowerPC Processing Element (PPE)
 - > General Purpose, 2-way SMT, SIMD instructions
 - > 8 Synergistic Processing Elements (SPE)
 - > Single threaded, specialized instruction set
 - > No direct main memory access, local storage instead (256KiB)
 - > 128bit SIMD instructions

- > Intel Xeon Phi
 - > 60 cores, 64 bit in order x86 architecture, up to 8 GiB RAM
 - > 4-way SMT (hide memory latency), 512bit SIMD instructions
 - > Behaves like a “normal” system, delivered as PCIe card

GPUs

- > Optimized for SIMD-like operations
 - > Many execution units, less resources for everything else
- > Optimized for throughput
 - > Single thread performance is not considered
 - > Latencies are hidden by multi-threading
- > E. g. NVIDIA GT200 family
 - > Up to 30 cores (“Streaming multiprocessors”, SMs)
 - > Each SM
 - > Can handle up to 32 threads groups (“warps”) of 32 threads each
 - > Has eight 32bit ALUs (i. e., a warp is executed in 4 steps)
 - > Has only one frontend (all threads in a warp must execute the same instruction)
 - > Has 16K registers, 16KiB shared memory

NVIDIA GT200




Source: realworldtech.com

General Purpose GPU Computing

- > Current vendor specific APIs
 - > *CUDA* for NVIDIA's GPUs
 - > *ATI Stream SDK* for AMD/ATI GPUs

- > Cross-vendor interface
 - > OpenCL (Open Computing Language) Specification
 - > Not only for GPUs
 - > Tries to make coprocessor processing power accessible
 - > E. g. also the "other" cores in a multi-core system
 - > OpenCL 1.0 published in 2008, now at OpenCL 1.2

OpenCL – Platform Model



- > *Compute devices (CDs)* are connected to a *host*
- > Compute devices consist of *compute units (CUs)* with *processing elements (PEs)*
- > PEs within a CU might execute in lockstep (SIMD)
- > CDs, CUs, and PEs have each individual memory
 - > A PE can access its own memory, and the memory of its CU and CD
 - > The host can only access the CD memories

OpenCL – Execution Model

- > The *Host program* submits operations to command queues, which are processed asynchronously
 - > Transfer memory to/from compute devices
 - > Execute *kernels* on compute devices
 - > (and synchronization commands)

- > A kernel is executed on a N-dimensional index space
 - > One kernel instance per index (*work-item*)
 - > Index space is decomposed into homogeneous *work-groups* (which have limited size)
 - > All work-items within a work-group are processed concurrently by PEs of one CU
 - > Cross CU synchronization/communication is not possible!

OpenCL – Kernels

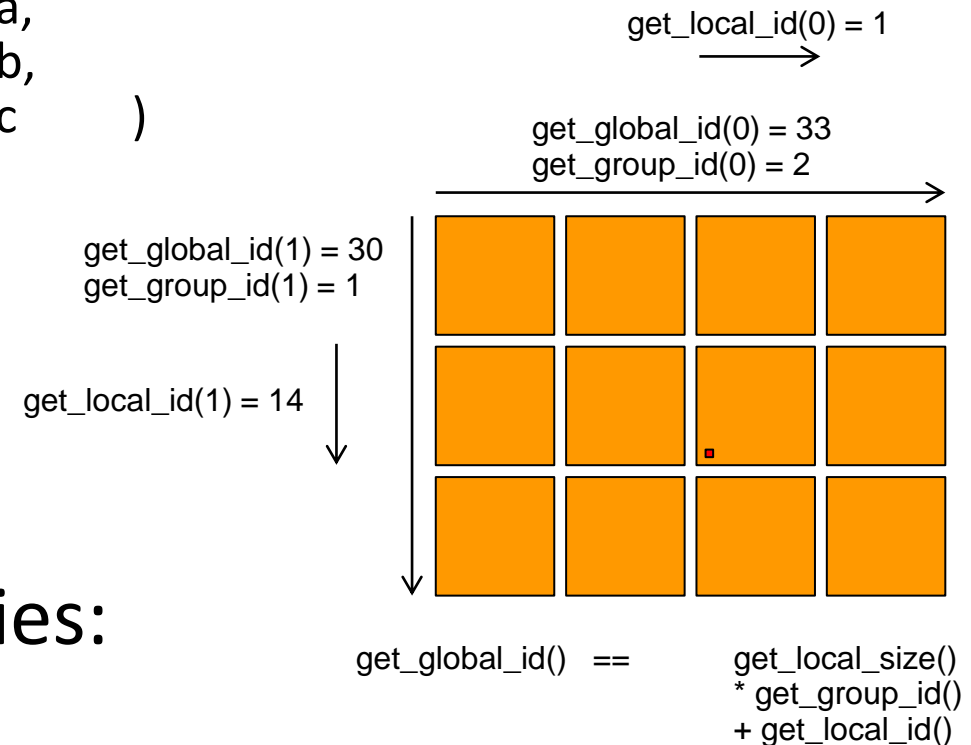
- > Kernels are (usually) written in OpenCL C
 - > Based on C99
 - > Some extensions (e. g. vector data types)
 - > Some restrictions (e. g. no recursion)

- > Kernels are (usually) included in source form with the (binary) host program
 - > OpenCL run-time includes a compiler
 - > Compiled explicitly during execution
 - > Code optimized towards the specific compute device
 - > (Caching of created binary is possible)

OpenCL C

> A simple OpenCL kernel

```
__kernel void vecadd(    __global float *a,
                        __global float *b,
                        __global float *c    )
{
    int x = get_global_id(0);
    c[x] = a[x] + b[x];
}
```



> Access to different memories:

- > Global memory: `__global`
- > Work-group memory: `__local`
- > Work-item memory: `__private` (or without qualifier)

Considerations

- > Work-items may execute in lockstep with some other work-items
 - > No synchronization necessary
 - > Code paths should not diverge

- > Work-items within a work-group are executed concurrently
 - > Synchronization possible (e. g., barrier)
 - > Should take advantage of local memory
 - > A variable declared `__local` is shared between all work-items in a work-group

- > The number of work-items in a work-group is limited
 - > Scalability is achieved by having many, many workgroups and a reasonable number of work-items per work-group
 - > Cross work-group synchronization is not possible
 - > (Only indirectly via the host; the actions of the current kernel are visible to the next kernel)

Control flow within a (simple) OpenCL host program



- > Query OpenCL run-time to find a suitable compute device
 - > E. g. `clGetPlatformIDs()`, `clGetDeviceIDs()`
- > Setup a context with associated devices
 - > E. g. `clCreateContext()`
- > Setup other things and associate them with the context
 - > Program objects: e. g. `clCreateProgramWithSource()`, `clBuildProgram()`
 - > Buffer objects: e. g. `clCreateBuffer()`
 - > Command queues: e. g. `clCreateCommandQueue()`
 - > Kernels: e. g. `clCreateKernel()`
- > Execute kernels
 - > Copy memory between host and device: e. g. `clEnqueueReadBuffer()`, `clEnqueueWriteBuffer()`
 - > Set kernel arguments: e. g. `clSetKernelArg()`
 - > Enqueue kernel: e. g. `clEnqueueNDRangeKernel()`
- > Free everything: `clRelease*()`