

# Parallel Programming

## OpenMP

Jan Schönherr

Technische Universität Berlin

School IV – Electrical Engineering and Computer Sciences

Communication and Operating Systems (KBS)

Einsteinufer 17, Sekr. EN6, 10587 Berlin

# OpenMP API

- > API for programming shared memory architectures
  - > Specified for C, C++ and Fortran
- > Contents
  - > Compiler directives
  - > Library routines (e. g. locks, query/set defaults)
  - > Environment variables (to change some defaults)
- > Features
  - > Partial/incremental parallelization of (existing) programs
  - > Portable source code
  - > Source code still compiles without OpenMP support

# OpenMP History

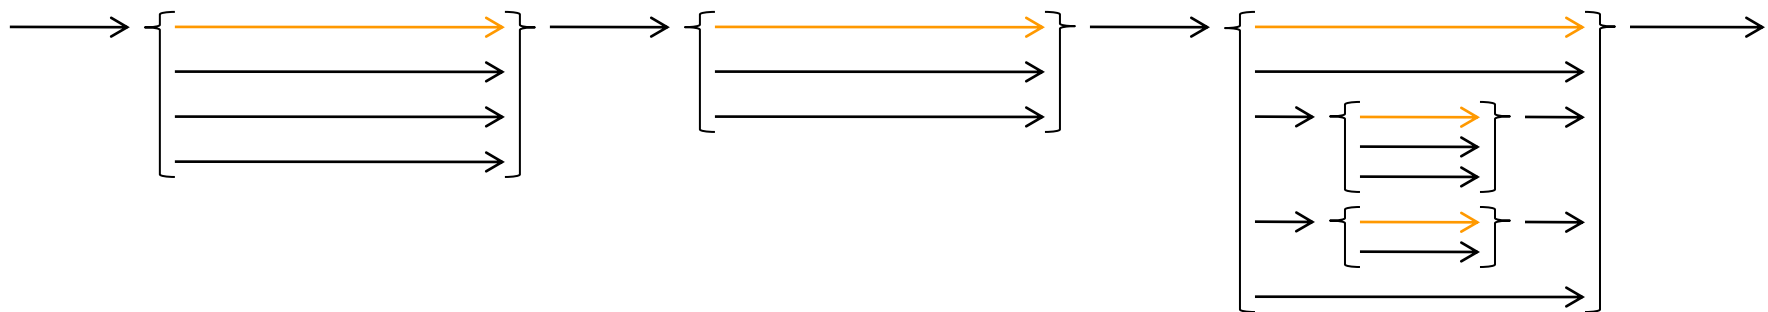
- > Early history
  - > OpenMP 1.0 in 1997 (Fortran) and 1998 (C/C++)
  - > OpenMP 2.0 in 2000 (Fortran) and 2002 (C/C++)
  - > OpenMP 2.5 in 2005
  
- > OpenMP 3.0 in 2008
  - > Adds *tasks* construct
  - > Supported with GCC 4.4, ICC 11.0
  - > Currently at OpenMP 3.1 in 2011
  
- > OpenMP 4.0 (currently RC2)
  - > Adds support for SIMD, coprocessors

# General C/C++ Syntax

- > Directives look like
  - > `#pragma omp <directive> [ clause, ... ]`
- > Most directives affect the next structured block
  - > One statement
  - > One compound statement (e. g. multiple statements enclosed in `{}`)
  - > Another OpenMP directive
- > Library functions are defined in `omp.h`

# Execution Model

- > OpenMP is based on Fork/Join parallelism
  - > Program execution starts with a single *initial thread*
  - > When told so, a thread creates a *team* of itself (the *master*) and zero or more additional threads
  - > Later, the team is joined
  - > In between is a *parallel region*



# Parallel Region

- > The *parallel* directive defines a parallel region
  - > A team of threads is created
  - > Each thread executes the structured block after the directive
  - > Only the master continues execution after the structured block
  - > Implicit barrier at the end
- > Number of threads depends on several factors
  - > E. g. your wishes, run-time system, implementation, ...
  - > It is fixed during the parallel region
- > Nesting of parallel regions possible, but potentially problematic
  - > Load balancing and overhead might become a problem
  - > Implementations may not support arbitrary nesting
    - > In that case, no additional threads are created

# Worksharing Constructs

- > Used inside of parallel regions
- > Work is shared among a team
  - > i. e. not all members execute everything inside the construct
- > Collective character, all members or none must reach it
- > Four constructs
  - > *for*: Divides loop iterations between members
  - > *sections*: Several structured blocks are distributed among members and executed in parallel
  - > *single*: A structured block is executed only by one member
  - > *workshare* (Fortran only)
- > Implicit barrier at the end (can be switched off)

# Parallel Loops

- > Iterations of a *for*-loop are divided among the team
- > Some restrictions regarding the loop
- > Different schedules possible
  - > *static*: equal-sized chunks are assigned round-robin
  - > *dynamic*: equal-sized chunks are assigned to free threads
  - > *guided*: chunks are assigned to free threads, starting large and becoming smaller
- > Selection of schedule may be delegated to the compiler and/or runtime system



# Parallel Statements

- > Execute (independent) statements in parallel
- > A *sections* construct is divided into multiple *sections*
  - > Sections may be executed concurrently, in any order
  - > Each section is executed by only one team member
  - > Distribution of sections among threads is implementation specific
- > Enables static distribution of work

# Tasks (since OpenMP 3.0)

- > Generalization of Sections, much more flexible
- > (Implicit) Shared task queue
  - > Every team member can create tasks
  - > Every team member can execute tasks
- > Dynamic work distribution
  - > Degree of parallelism limited by team size
- > Limited support for synchronization between tasks
  - > Team members may switch between multiple tasks
  - > Task may yield execution explicitly
  - > Task can wait for completion of all child tasks
- > For, e. g., divide & conquer and other cases of nested parallelism

# Data-sharing (i)

- > Data-sharing clauses specify what happens to variables before, during and after an OpenMP construct
- > (Not all clauses are accepted on all directives)
- > Variables declared before a construct can be set to:
  - > *shared*: all threads access the same memory
  - > *private*: each thread accesses its own uninitialized memory
  - > *firstprivate*: like *private*, but initialized
  - > *reduction*: each thread gets a local, appropriately initialized copy; origin is updated afterwards with the result of a reduction
  - > *copyprivate*: update all private origins after a single construct
  - > ... (some more for more specialized uses)

# Data-sharing (ii)

- > Data-sharing defaults
  - > Parallel region: *shared*, with some exceptions
    - > e. g. a parallel loop variable is *private*
  - > Tasks: mostly *firstprivate*
- > Default can be disabled with *default(none)*
  - > Prevents errors as you must now specify each attribute explicitly
- > Data-sharing attributes on pointers/arrays might not do what you think

# Synchronization Constructs

- > *barrier*: Execution continues only after all members have reached the barrier
- > *master*: Block is only executed by the master of the team
- > *critical*: Block is executed by at most one thread at the same time
- > *atomic*: A memory location is updated atomically (and since OpenMP 3.1 optionally returned)
- > *flush*: Synchronizes a thread's view of shared memory

# Behind the scenes of GNU OpenMP

```
#pragma omp parallel
{
    body;
}
```



```
void subfunction (void *data)
{
    use data;
    body;
}

setup data;
GOMP_parallel_start (subfunction, &data, num_threads);
subfunction (&data);
GOMP_parallel_end ();
```

```
#pragma omp for schedule(runtime)
for (i = 0; i < n; i++)
    body;
```



```
{
    long i, _s0, _e0;
    if (GOMP_loop_runtime_start (0, n, 1, &_s0, &_e0))
        do {
            for (i = _s0, i < _e0; i++)
                body;
        } while (GOMP_loop_runtime_next (&_s0, &_e0));
    GOMP_loop_end ();
}
```